



# Lecture VI

## A DevOps Journey : Container Basics

From Virtual Machines to Containers  
Understanding Container Technology Fundamentals  
Guy Van Eeckhout  
2025-2026

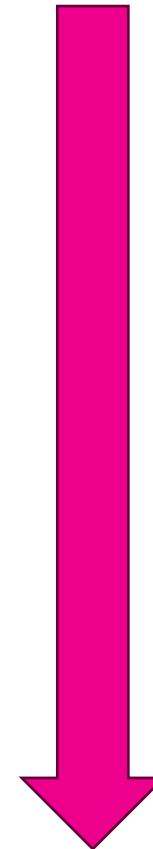


(inspired by Alex Desmedt)

# What we'll talk about

---

- The historical evolution of container technology (1979-2013)
- How Docker emerged and revolutionized containerization
- Key milestones in Docker's evolution (2013-2025)
- How Docker positions itself against alternatives (Podman, Kubernetes)
- The concept of Docker registries and repositories



# CONTAINERS VS VIRTUAL MACHINES (yay another recap)

---

## Fun with VMWARE (until now)

- Multiple VMs on one physical machine
- Each VM: Full OS + Applications
- Hypervisor manages hardware access
- Heavy resource consumption

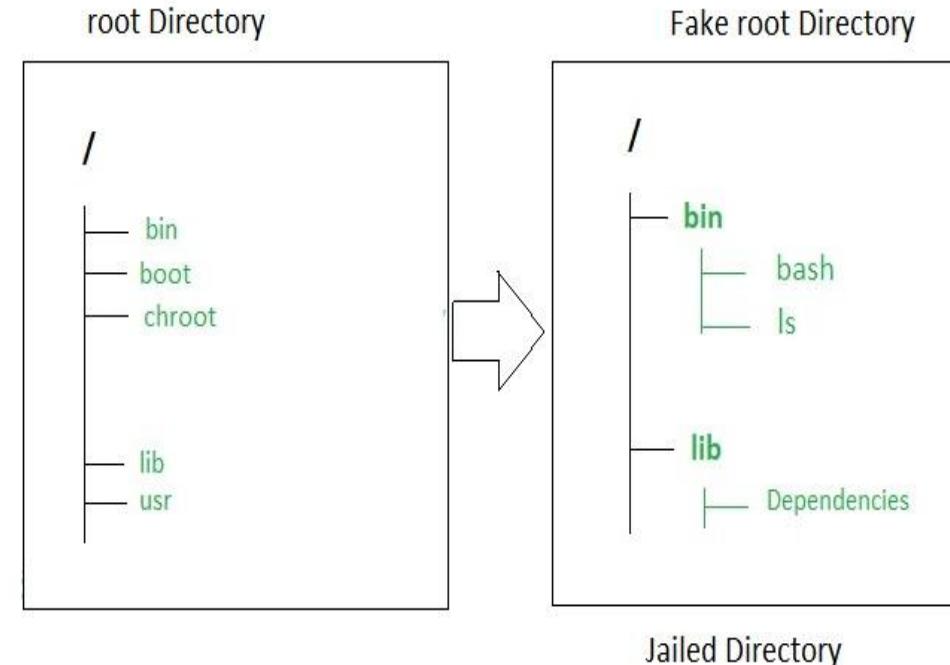
## The container approach

- Multiple containers on one OS
- Shared kernel, isolated processes
- Container runtime manages isolation
- Lightweight and efficient

In short: Containers virtualize the OS, VM's virtualize the hardware

# Humble beginnings : 1979

- UNIX v7 introduces CHROOT
  - *Change root* command
  - Changes root directory of a process
  - First step toward process isolation
- What was the purpose of chroot ?
  - Originally used for testing new environments
  - Isolated filesystem access for processes
  - Created "chroot jail" concept
- Limitations
  - Only filesystem isolation
  - Processes could still communicate
  - Not suitable for security isolation alone



Still available on current linux systems !

# 21 years later : 2000

- FreeBSD Jails (2000) - Beyond Filesystem Isolation
  - Building further on chroot:
  - Developed by FreeBSD for hosting providers
  - Addressed security and administration needs
  - Partitioned FreeBSD system into independent "jails"

#### **Enhanced Features:**

- ✓ Own IP address per jail
  - ✓ Own filesystem
  - ✓ Own configuration
  - ✓ Better process isolation
  - ✓ Independent user spaces

## Use Case:

- Hosting providers separating customer services
  - Security enhancement over chroot
  - True multi-tenant environment



Still available on current FreeBSD systems !

# 2001 : Vserver : Linux enters the fray

---

Parallel Development:

- Similar to FreeBSD Jails, but for Linux
- Operating system virtualization via kernel patches
- Created "Security Contexts" and VPS (virtual private server)

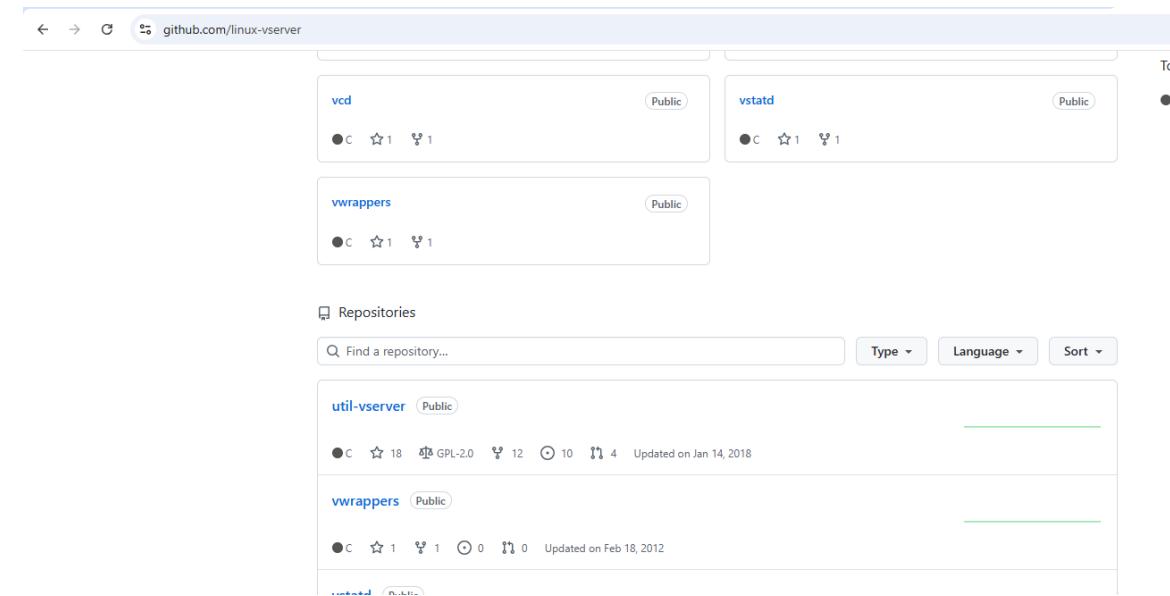


# Vserver : Linux enters the fray

Isolated Resources:

- File systems
- Network addresses
- Memory allocation
- CPU time
- Process trees

Paved way for future Linux containers



Status: Last stable release: 2006

# 2004 – 2005 : Solaris Zones & OpenVZ

---

## Solaris Zones (2004):

Introduced in Solaris 10 (Sun Microsystems/Oracle)

Lightweight virtualization for SPARC and x86

Combined resource controls + boundary separation

Two modes: Container mode & Whole system mode

## OpenVZ (2005):

Developed by Swsoft

Patched Linux kernel approach

Similar features to Solaris Containers

Marked kernel-level virtualization was going mainstream

Impact: Enterprise adoption begins. Proven production-ready technology. Foundation for future container tech



# 2006 PROCESS CONTAINERS (CGROUPS)

---

Process Containers (2006)

- Later renamed "Control Groups" (cgroups)
- Merged into Linux kernel 2.6.24
- Limited, accounted, and isolated resource usage

cgroups control:

- CPU allocation
- Memory limits
- Disk I/O
- Network bandwidth



# 2008 : Linux Containers (LXC)

---

IBM development

BUT with mainstream Linux kernel integration

Combined cgroups + namespaces

First complete container solution

Used kernel features: namespaces, cgroups, chroots, capabilities

Foundation Complete: All pieces in place for modern containers!



# And now Docker

So turn off your lights

# Reality Check from last week: No Docker vs Docker

---

## No Docker

- ✗ "Works on my machine" syndrome
- ✗ Complex dependency management
- ✗ Different dev/test/production environments
- ✗ Difficult application deployment
- ✗ Lengthy setup times
- ✗ Environment configuration drift

## Docker

- ✓ Package application + dependencies together
- ✓ Consistent behavior across environments
- ✓ Simple deployment process
- ✓ Fast startup times
- ✓ Easy sharing and distribution
- ✓ Developer-friendly tools

Solomon Hykes' Vision:"Shipping code to the server is hard" - make it easy!

# 2008-2010 : PLATFORM AS A SERVICE

---

Where Docker was born

Company Formation:

Founded 2008 in Paris

by Kamel Founadi, Solomon Hykes, Sebastien Pahl

Incorporated in US in 2010

Joined Y Combinator Summer 2010

The Vision:

- Platform-as-a-Service (PaaS) provider
- Enable easy cloud deployment
- Manage applications in the cloud

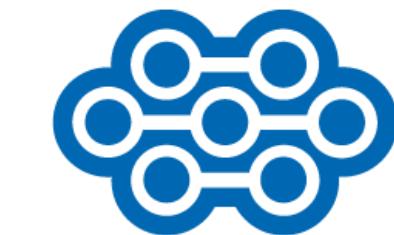


# 2008-2010 : PLATFORM AS A SERVICE

---

From an Internal Project:

- Docker started as internal tool (2010-2012)
- Solved dotCloud's own deployment challenges
- Used container technology to isolate processes
- Improved resource management for PaaSKey



**dotCloud**



The tool they built was more valuable than the service!

# PyCon March 13, 2013 : The day everything changed



The Big Reveal: Solomon Hykes demonstrates Docker publicly

5-minute demo that changed software development

The Demo Message: "For developers, shipping to the server is hard"

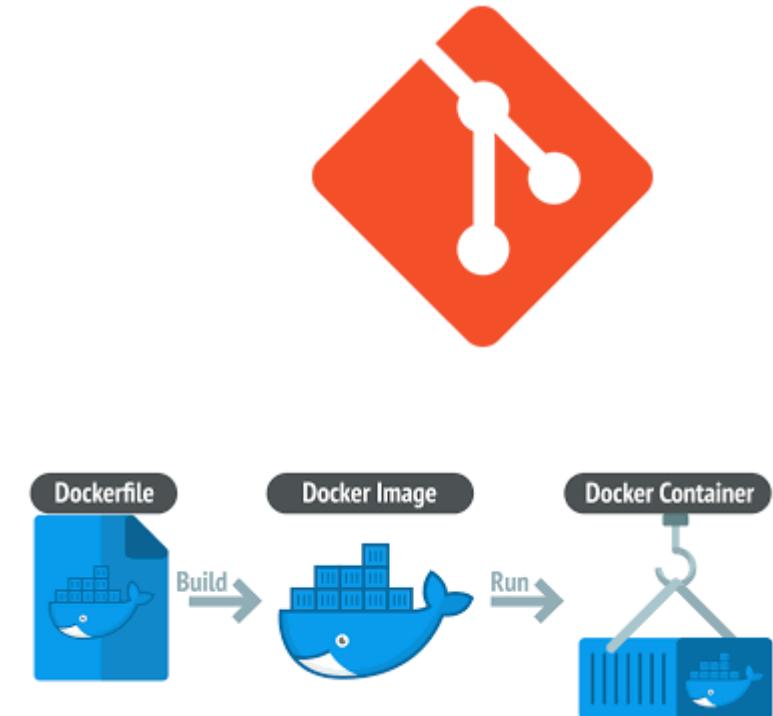
A composite image showing Solomon Hykes speaking at a podium with a Docker logo, and a terminal window displaying Docker logs. The terminal output shows multiple container instances starting and exiting, with the message "echo hello world" visible in one of the logs.

```
ago  Exit -1
5d117b05  c5860e37a1bcb678  /bin/sh
ago  Exit 0
7bbf2ed7  fad7c890e7ca7be4  /bin/bash
ago  Exit 0
1c57729e  c5860e37a1bcb678  /bin/echo hello worl About an hour
ago  Exit 0
867600d5  fad7c890e7ca7be4  /bin/bash
ago  Exit -1
0cd7b00d  c5860e37a1bcb678  /bin/sh
ago  Exit 0
195e0591  c5860e37a1bcb678  /b
ago  Exit -1
1a832a1e  c5860e37a1bcb678  /b
ago  Exit -1
40d49a1a  c5860e37a1bcb678  /b
ago  Exit 0
f318a96d  c5860e37a1bcb678  /b
ago  Exit -1
8c0f9431  c5860e37a1bcb678  /b
ago  Exit 0
$
```

# Why Docker was revolutionary part 1

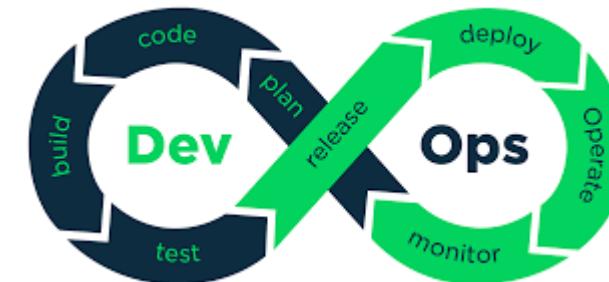
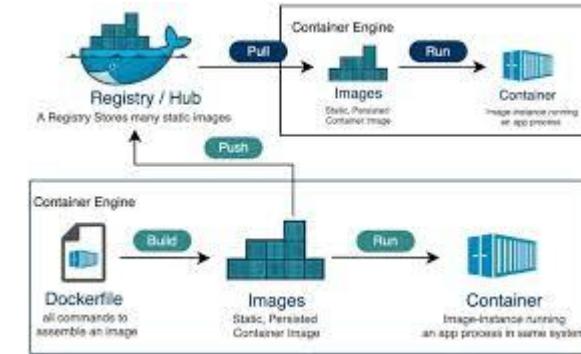
---

- Developer Experience:
  - Git-like semantics for containers
  - Simple CLI commands
  - Dockerfile for reproducible builds
- Standardized Format:
  - Universal container image format
  - "Build once, run anywhere"
  - Separates app from infrastructure



# Why Docker was revolutionary part 2

- Ecosystem Integration:
  - Docker Hub for sharing images
  - Integration with existing tools
  - Rich community and marketplace
- Bridged Dev and Ops:
  - Developers focus on "inside container"
  - Operations focus on "outside container"
  - Reduced friction between teams



# Docker Early Architecture : LXC Based

---

## Initial Implementation:

- Built on top of LXC (Linux Containers)
- LXC as default execution environment
- Docker added user-friendly layer
- Simplified container management

## Limitations of LXC Dependency:

- Tied to Linux-specific implementation
- Limited portability
- Dependent on external project
- Needed more control

## Architecture Components:

Docker CLI

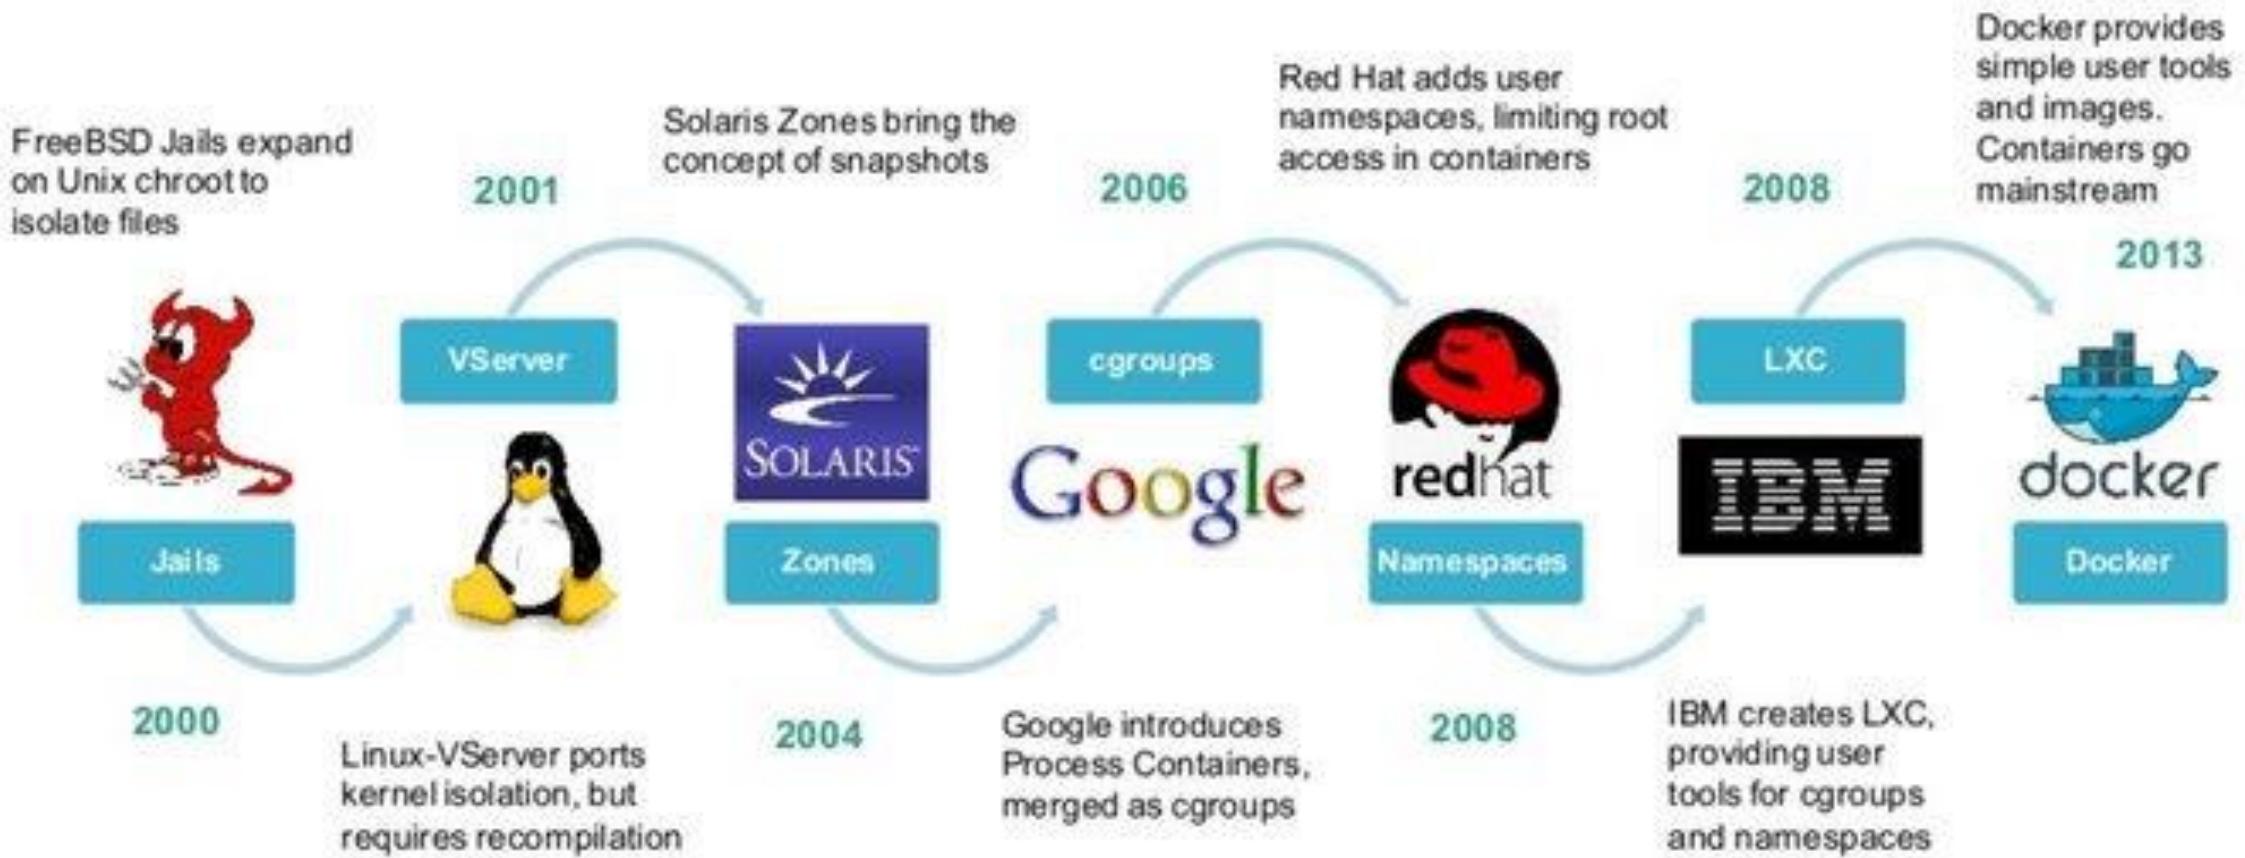
Docker Daemon

LXC

Linux

This led to Docker's biggest technical evolution...

# We got to here : (recap)



# DOCKER'S INNOVATION - SIMPLIFYING CONTAINERIZATION

---

Before Docker:

LXC existed but was complex

Required deep Linux knowledge

Manual configuration required

Steep learning curve



Docker's Abstraction:

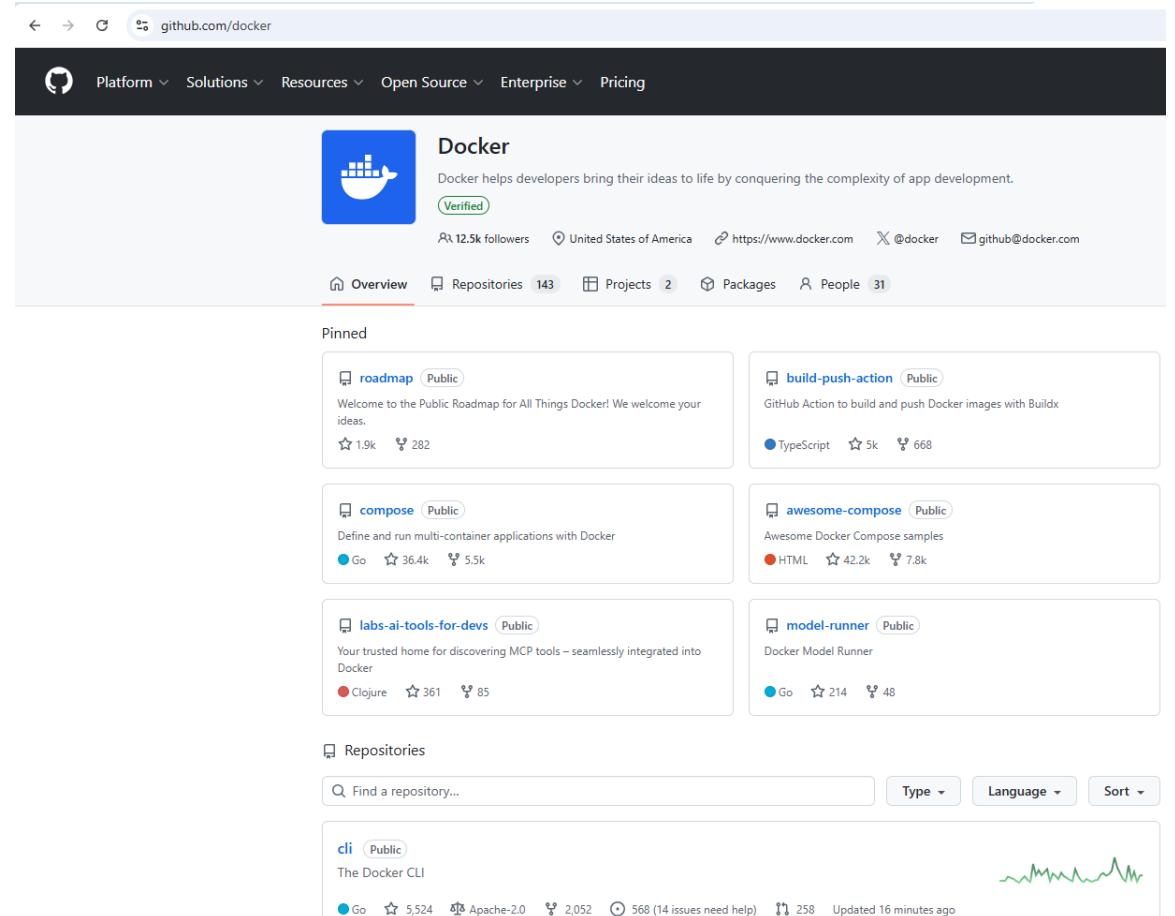
1. Dockerfile - Simple text file defines container
2. Docker Images - Layered, reusable, cacheable
3. Docker Hub - Central repository for sharing
4. Docker CLI - Intuitive commands
5. Docker Compose - Multi-container applications

`docker run nginx`

# 2013 : Docker goes open source

## The Release:

- Code made publicly available March 2013
- Apache 2.0 License
- GitHub repository created
- Community development begins
- Impact of Open Source:
  - ✓ Rapid community adoption
  - ✓ Community contributions and feedback
  - ✓ Ecosystem development
  - ✓ Trust and transparency
  - ✓ Innovation acceleration



# 2013 : Docker goes open source

---

By The Numbers (First Year):

- Thousands of contributors
- Rapid feature development
- Growing ecosystem
- Industry attention



Strategic Wisdom: Open sourcing Docker was key to its success

Enabled widespread adoption and innovation far beyond what a closed-source product could achieve.

# 2014 - FROM LXC TO LIBCONTAINER

---

The Transition:

- Docker 0.9 release (March 2014)
- Introduced libcontainer (Written in Go programming language)
- Replaced LXC as default execution driver
- Why libcontainer?
  - ✓ Docker-owned and maintained
  - ✓ Cross-platform potential
  - ✓ Direct access to kernel features
  - ✓ Better performance
  - ✓ More control over functionality

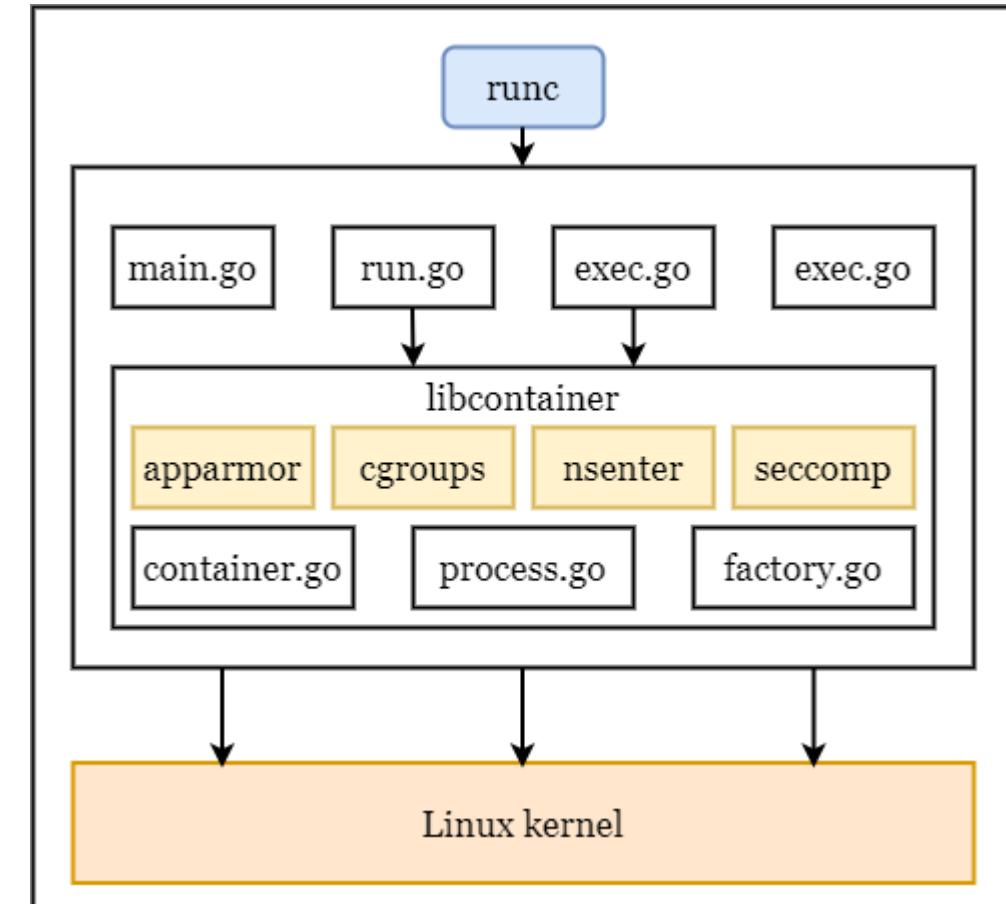


# 2014 – LIBCONTAINER (evolved into => runc)

Technical Details:

- Direct interface to:
  - Linux namespaces
  - Control groups (cgroups)
  - Capabilities
  - AppArmor/SELinux profiles

Result: Docker became independent and portable!

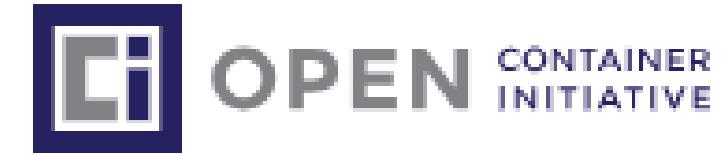


# 2015 - THE OPEN CONTAINER INITIATIVE (OCI)

---

What is OCI?

- Founded June 2015
- Hosted by Linux Foundation
- Open governance structure
- Created by Docker, CoreOS, and others



Wat is its Mission ?

- Establish industry standards for containers:
  1. Runtime Specification - How containers execute
  2. Image Specification - Container image format

# 2015 - THE OPEN CONTAINER INITIATIVE (OCI)

---

Docker's Contribution:

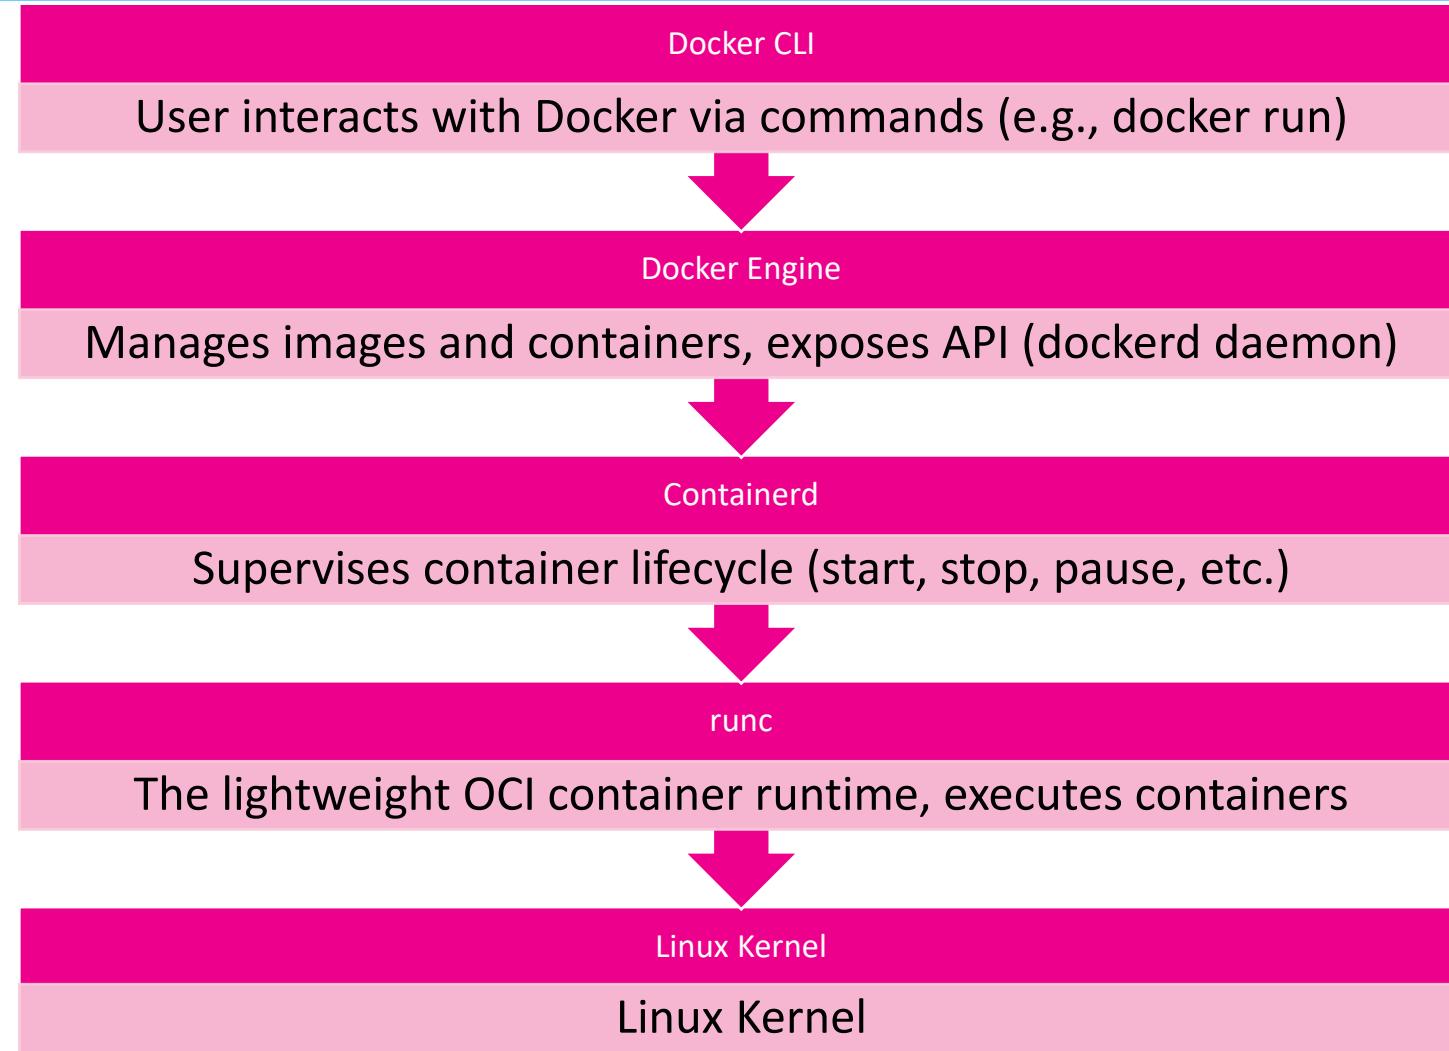
- Donated runc (formerly libcontainer)
- Reference implementation of OCI runtime spec
- Later donated containerd

Why Standards Matter:

- Prevents fragmentation
- Ensures interoperability
- Enables innovation
- Industry-wide adoption
- Vendor neutrality

Impact: Docker containers work everywhere!

# The current docker full stack



# What about docker desktop ?

Layer	Role/Contribution
Docker Desktop	<p>The orchestrator and facilitator for Docker on desktops</p> <ul style="list-style-type: none"><li>- Provisioning VM: Runs a lightweight Linux VM using Hyper-V, WSL2, or Apple Hypervisor for kernel features on non-Linux OSes.</li></ul>
	<ul style="list-style-type: none"><li>- UI and Integration: Offers graphical dashboards, settings, extensions, and status monitors unavailable via plain Docker CLI.</li></ul>
	<ul style="list-style-type: none"><li>- Install/Run: Installs, launches, and updates Docker Engine and related binaries.</li></ul>
Docker CLI	User commands (text UI) get passed to the Docker Engine running inside the provisioned VM.
Docker Engine	Runs inside the VM, unchanged from how it would on a native Linux host.
containerd, runc, Linux Kernel	Operate inside the VM where Docker Engine runs, using Linux kernel features enabled by the virtual machine.

# Docker desktop & Docker Hub

---

Docker Desktop:

Launched for Mac and Windows

2022: Linux support added

Local development environment

Integrated Kubernetes

Developer-focused features

Docker Hub:

Central registry for container images

Public and private repositories

Official images maintained by Docker

Automated builds

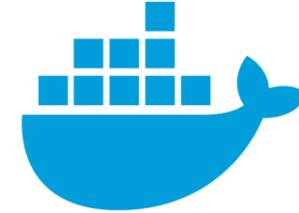
15+ million repositories

# The Current Container Landscape

---

Container Runtimes:

- Docker - Most popular, developer-friendly
- Podman - Daemonless alternative (Red Hat)
- containerd - CNCF graduated project
- CRI-O - Kubernetes-native runtime

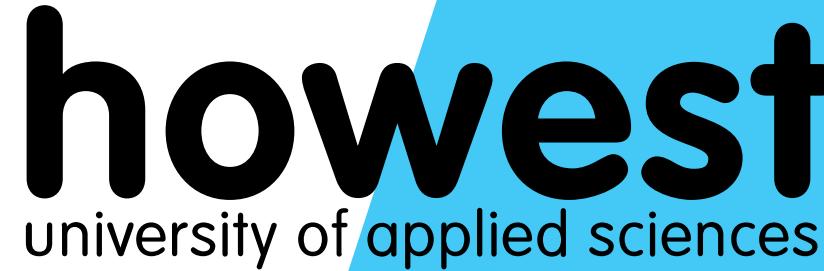


**podman**



**cri-o**

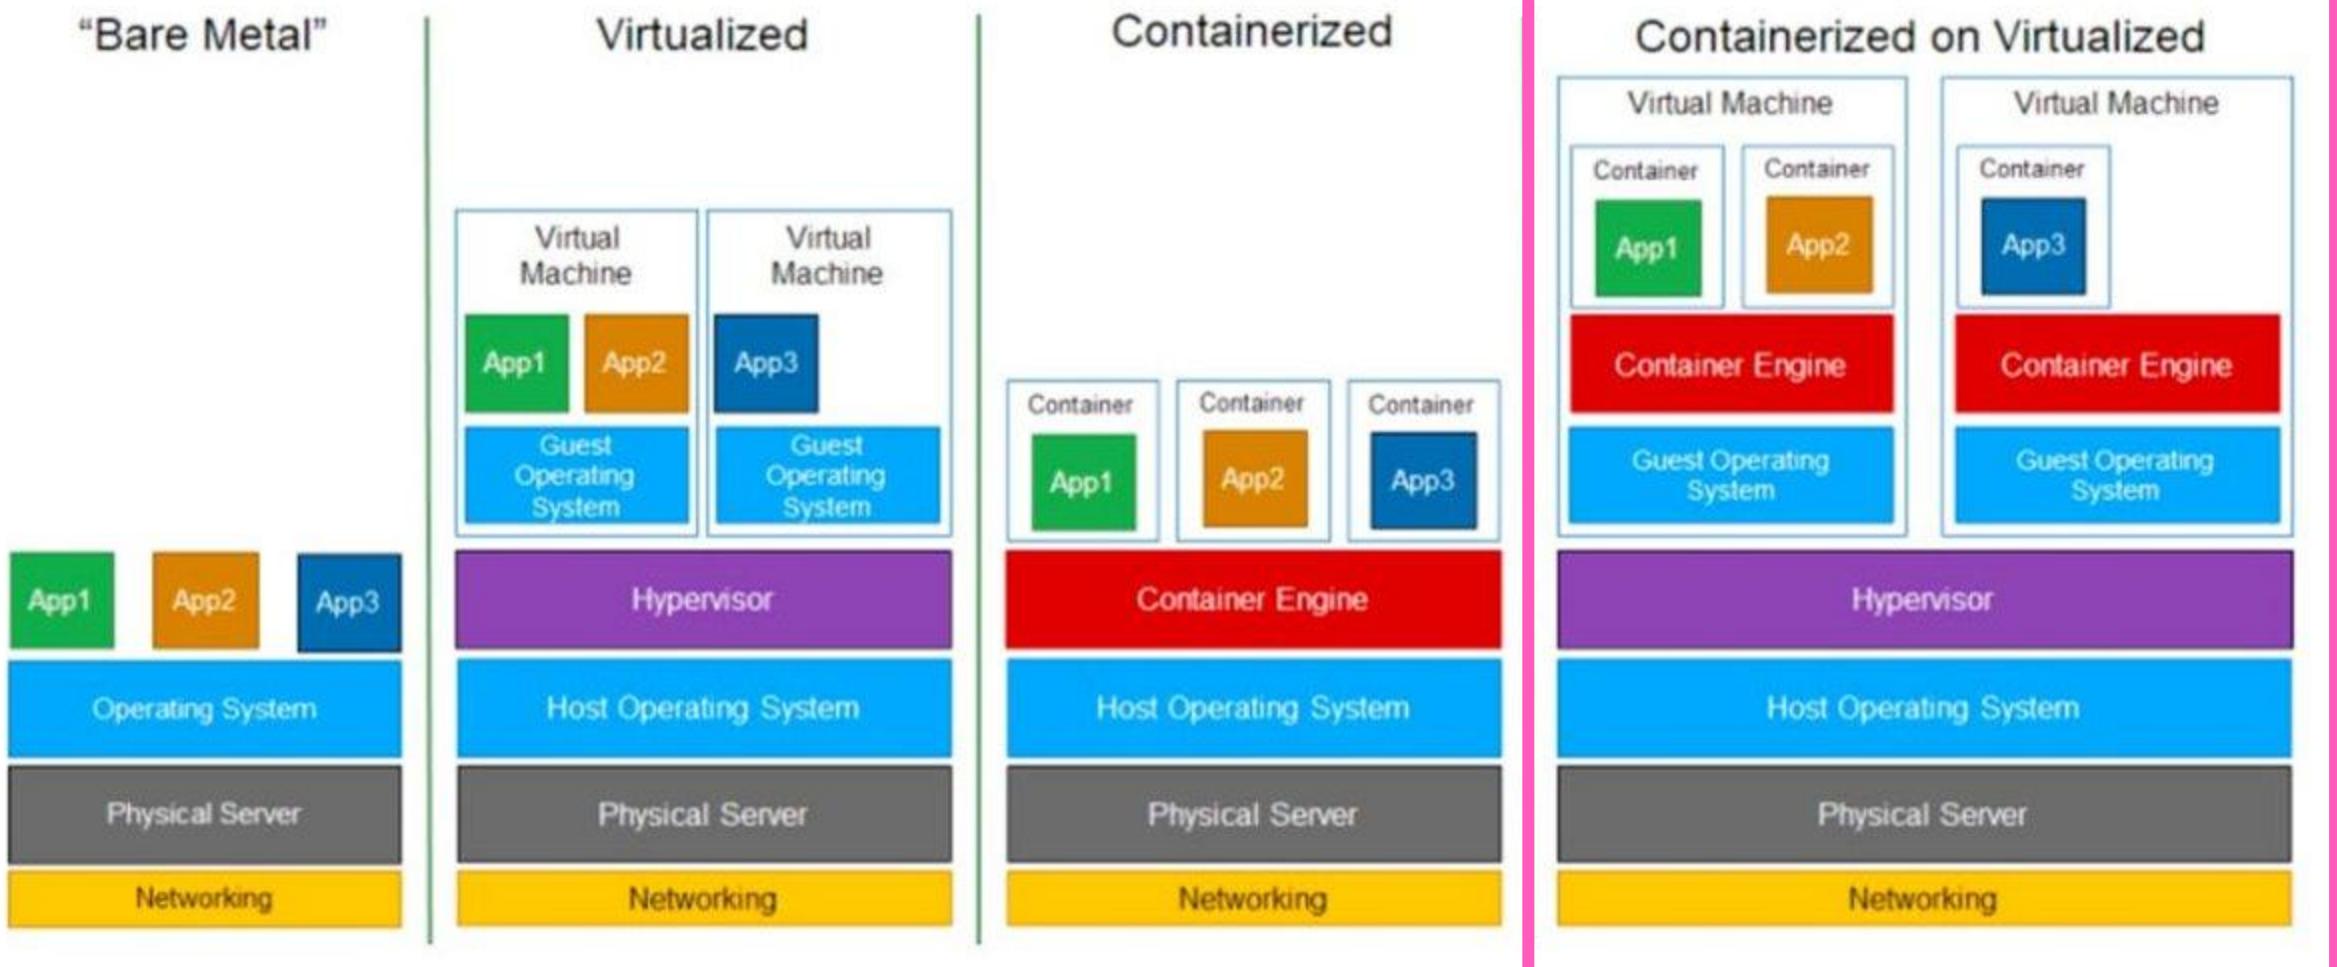
**containerd**



# Getting your hands dirty (finally)

With Docker ! 😊

# Which setup should you choose ?



# So sorry, we won't be using

The screenshot shows the Docker Desktop application interface. The left sidebar has a 'Containers' tab selected. The main area displays a search bar, a toggle for 'Only show running containers' (which is off), and a table of containers. The table has columns for Name, Container ID, Image, Port(s), CPU (%), Last st, and Actions. There are six entries in the table, all of which are stopped.

	Name	Container ID	Image	Port(s)	CPU (%)	Last st	Actions
<input type="checkbox"/>	lab6	01bf87661e8f	<a href="#">caoslab6</a>		N/A	18 hours	<a href="#">▶</a> <a href="#">⋮</a> <a href="#">trash</a>
<input type="checkbox"/>	elegant_tharp	0291553971c1	<a href="#">hello-world</a>		N/A	2 months	<a href="#">▶</a> <a href="#">⋮</a> <a href="#">trash</a>
<input type="checkbox"/>	sharp_kirch	1513c230c809	<a href="#">ubuntu</a>		N/A	1 month	<a href="#">▶</a> <a href="#">⋮</a> <a href="#">trash</a>
<input type="checkbox"/>	caoslab3	5cbd0aabda45	<a href="#">koenkorem</a>		N/A	1 month	<a href="#">▶</a> <a href="#">⋮</a> <a href="#">trash</a>
<input type="checkbox"/>	kasmcaos	7d40c877c636	<a href="#">koenkorem</a>	6901:6901	N/A	1 month	<a href="#">▶</a> <a href="#">⋮</a> <a href="#">trash</a>

# Problem 1 : installing docker

On Debian :

apt install docker(.io) ?

Will (usually work) BUT !

Version Lag: Debian Stable prioritizes stability over recency. Docker versions in Debian repositories can be months or even years behind current releases

Limited Updates: Debian Stable only receives security updates and critical bug fixes, not feature updates.

Incomplete Package Set: The default Debian packages may not include all the modern Docker components



# Solution 1 : installing docker the right way

Docker has its OWN package repository

But you need to tell Debian to trust that repository (it doesn't do so by default)



# Installing docker “the right way”

---

## 1. Install Prerequisites

```
sudo apt update
```

```
sudo apt install apt-transport-https ca-certificates curl gpg
```

## 2. Add Docker's GPG Key

```
curl -fsSL https://download.docker.com/linux/debian/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker.gpg
```

## 3. Configure the Repository (all in ONE line !)

```
echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker.gpg]  
https://download.docker.com/linux/debian trixie stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

# Installing docker “the right way”

---

## 4. Install Docker

*sudo apt update*

*sudo apt install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin*

## 5. Check if Docker is started and enabled, if not Enable and Start Docker system service

*systemctl status docker (is it running / enabled ? If not :)*

*sudo systemctl start docker*

*sudo systemctl enable docker*

## 6. Configure User Permissions for your regular user(to allow non-root users to run docker)

*sudo usermod -aG docker \${USER}*

*newgrp docker*

## 7. Check your installation

*docker --version*

# Docker commands you need to know blindly

Command	Purpose
<b>docker run</b>	Start a container image
<b>docker pull</b>	Download (pull) a container image, but do NOT create a container
<b>docker build</b>	Build your own image
<b>docker exec</b>	Execute a command inside of your running container.
<b>docker start/stop</b>	Start/stop a container
<b>docker inspect</b>	Retrieve all the metadata of a container.
<b>docker ps</b>	Get a list of all (running) containers
<b>docker logs</b>	Get the logs out of a container

# Example commands

---

**docker run hello-world**

**docker run -d httpd** (detached mode) ; **docker run -d fedora sleep infinity** (keeps alive)

**docker start/stop <container-id or name>**

**docker rm <container-id or name>**

**docker exec -ti <container-id or name> <command>**

e.g. :

```
root@mis1:~# docker exec -ti my_web /bin/bash
root@38ed783b3b1e:/usr/local/apache2# cat /etc/os-release
PRETTY_NAME="Debian GNU/Linux 13 (trixie)"
NAME="Debian GNU/Linux"
VERSION_ID="13"
VERSION="13 (trixie)"
VERSION_CODENAME=trixie
DEBIAN_VERSION_FULL=13.1
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
root@38ed783b3b1e:/usr/local/apache2# whoami
root
```

**shell in a container: docker exec -ti <container-id or name> /bin/bash**

# Intermission : TAG rant

Please do it right

# Container image tagging

---

Tag = metadata

Convenient for version control

Your choice (but please stick to common good practise)

Assigned when building a container image (e.g. **docker build -t <image\_name>:<image\_tag>**)

Or added later:

**docker image tag SOURCE\_IMAGE[:TAG] TARGET\_IMAGE[:TAG]**

# Container image tagging

---

- **latest**: special tag, used as default
- Can be text, digit or any combination

Tag	Image specifics
nginx:1.27.1	Exact version, no specification of OS.
nginx:mainline	No version tag, will update when tag is changed (not recommended)
nginx:1	Only specify major version. Every partial update will occur
nginx:1.27	Minor updates will occur.
nginx:1.27.1-trixie	Exact version pin + OS (Debian 13)
nginx:1.27.1-alpine	Exact version pin + OS (Alpine, no version)

# Container image tagging – best practices

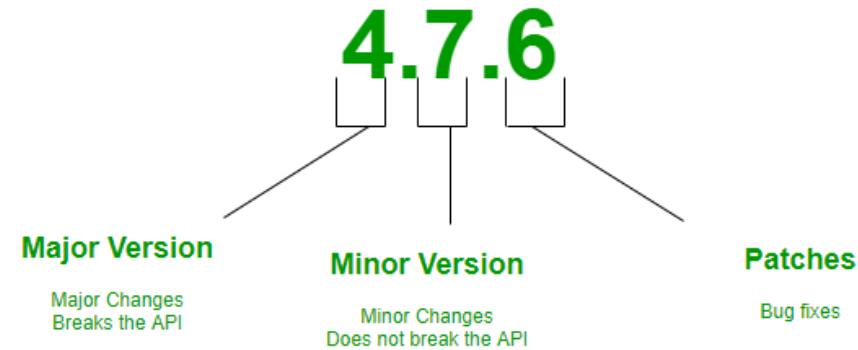
---

- Always add a **version pin** to your container images
- Never use latest, mainline or images without version info in production!
  - Automatic builds will update warning
  - This WILL break your pipeline. Trust me 😅
- Choose a standardized image tag across services and STICK WITH IT
  - **v1.3.26 != 1.3.26**

# Container image tagging – Semantic versioning

---

- Industry standard to tag container images
- MAJOR.MINOR.PATCH
- <https://semver.org/>
- Integrates great with git tagging !
  - `git tag 4.7.6 # tag this commit`
  - `git tag -a 1.0.0 -m "Release version 1.0.0"`
  - `git push origin 1.0.1`
  - `git push --tags`



**End of TAG rant**

Thank you for your attention

# Docker images

---

- Overview: **docker image ls**

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker image ls
REPOSITORY      TAG          IMAGE ID      CREATED       SIZE
httpd           latest        6a4fe18d08d2  4 days ago   117MB
hello-world     latest        1b44b5a3e06a  3 months ago  10.1kB
osc-guy-van-eec@lnx-guy-van-eec:~$
```

- (look at the SIZE)
- Images are stored at a *container registry*
  - hub.docker.com (Default Docker Registry)
  - quay.io (RedHat)
  - <https://images.linuxcontainers.org/> (LXC containers)

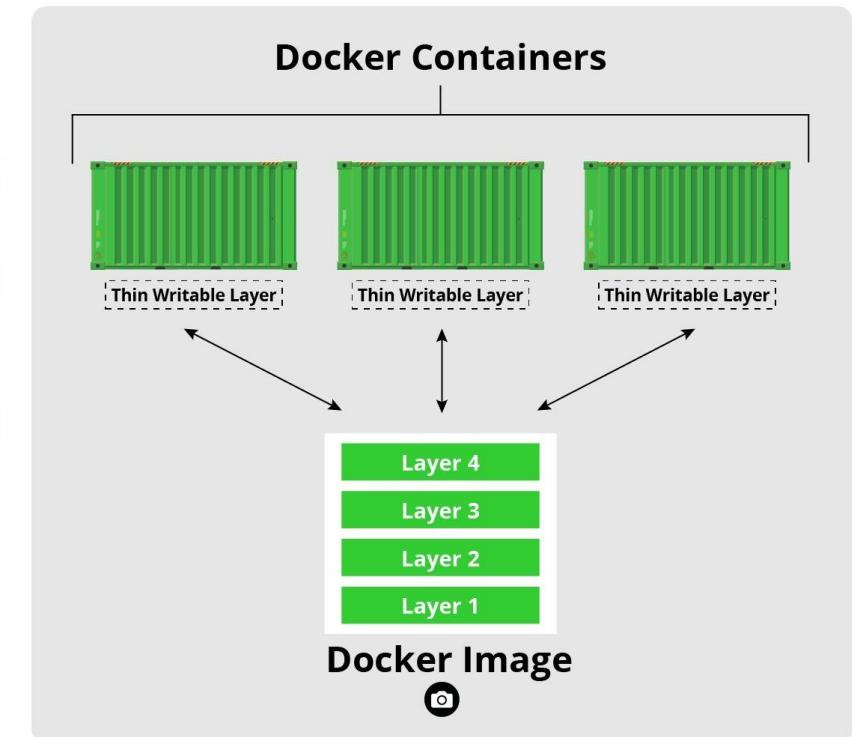
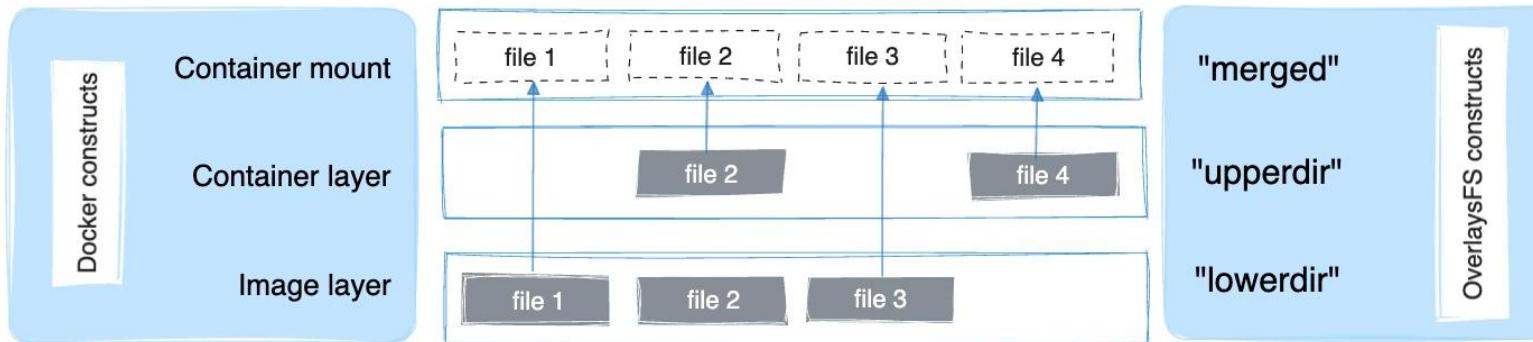
# Some interesting Docker images

---

- Base OS containers:
  - Fedora: [https://hub.docker.com/\\_/fedora](https://hub.docker.com/_/fedora)
  - Debian: [https://hub.docker.com/\\_/debian](https://hub.docker.com/_/debian)
  - Ubuntu: [https://hub.docker.com/\\_/ubuntu](https://hub.docker.com/_/ubuntu)
- Preinstalled Services:
  - Httpd (Apache): [https://hub.docker.com/\\_/httpd](https://hub.docker.com/_/httpd)
  - Nginx: [https://hub.docker.com/\\_/nginx](https://hub.docker.com/_/nginx)
  - Locust (load testing): <https://hub.docker.com/r/locustio/locust>
- Even full applications:
  - Wordpress: <https://hub.docker.com/u/library>
  - Netbox: <https://hub.docker.com/r/netboxcommunity/netbox>

# Docker images

- Consists of layers. Only the top layer is writable !
- Overlay2 fs (/var/lib/docker/overlay2/)



- Layers that are shared between containers are only stored ONCE

# Docker images

---

- Remove images: **docker image rm <image-naam>**

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker image rm hello-world
Untagged: hello-world:latest
Untagged: hello-world@sha256:56433a6be3fda188089fb548eae3d91df3ed0d6589f7c2656121b911198df065
Deleted: sha256:1b44b5a3e06a9aae883e7bf25e45c100be0bb81a0e01b32de604f3ac44711634
Deleted: sha256:53d204b3dc5ddbc129df4ce71996b8168711e211274c785de5e0d4eb68ec3851
osc-guy-van-eec@lnx-guy-van-eec:~$
```

You cannot remove an image that is still being used by a container ! (without force)

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker image rm hello-world
Error response from daemon: conflict: unable to remove repository reference "hello-world" (must force) - container d77a0d37a812
is using its referenced image 1b44b5a3e06a
osc-guy-van-eec@lnx-guy-van-eec:~$ docker image rm --force hello-world
Untagged: hello-world:latest
Untagged: hello-world@sha256:56433a6be3fda188089fb548eae3d91df3ed0d6589f7c2656121b911198df065
Deleted: sha256:1b44b5a3e06a9aae883e7bf25e45c100be0bb81a0e01b32de604f3ac44711634
```

- Remove all unused images **docker image prune**
  - Only if images have not been used for a while (clean up / regain discspace)

# Docker networking

---

- All containers have default network access
- 3 default networks available at install (**docker network ls**):

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker network ls
NETWORK ID      NAME      DRIVER      SCOPE
6c7f97ed6a31   bridge    bridge      local
706f7d0c3b75   host      host       local
fa28d1a302fa   none      null       local
osc-guy-van-eec@lnx-guy-van-eec:~$ |
```

# Docker networking

- Find out more info about these networks
  - **docker network inspect <network>**

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker network inspect bridge
[{"Name": "bridge",
 "Id": "6c7f97ed6a31811c3e5cecc35a30123d26deba0bb79d4ec512da26935d02c420",
 "Created": "2025-11-08T10:07:05.10610615+01:00",
 "Scope": "local",
 "Driver": "bridge",
 "EnableIPv4": true,
 "EnableIPv6": false,
 "IPAM": {
     "Driver": "default",
     "Options": null,
     "Config": [
         {
             "Subnet": "172.17.0.0/16",
             "Gateway": "172.17.0.1"
         }
     ]
 },
 "Internal": false,
 "Attachable": false,
 "Ingress": false,
 "ConfigFrom": {
     "Network": ""
 },
 "ConfigOnly": false,
 "Containers": {},
 "Options": {
     "com.docker.network.bridge.default_bridge": "true",
     "com.docker.network.bridge.enable_icc": "true",
     "com.docker.network.bridge.enable_ip_masquerade": "true",
     "com.docker.network.bridge.host_binding_ipv4": "0.0.0.0",
     "com.docker.network.bridge.name": "docker0",
     "com.docker.network.driver.mtu": "1500"
 },
 "Labels": {}
}]
```

# Docker networking

- Request Network info of a container
- In container itself (**ip a** or **ifconfig**)
- **docker inspect <container-id>**

```
[root@5e274cbcad4e /]# ip a s
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host proto kernel_lo
            valid_lft forever preferred_lft forever
2: eth0@if14: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether a2:a7:fc:ed:cd:40 brd ff:ff:ff:ff:ff:ff link-netnsid 0
        inet 172.17.0.2/16 brd 172.17.255.255 scope global eth0
            valid_lft forever preferred_lft forever
[root@5e274cbcad4e /]# |
```

```
"HairpinMode": false,
"LinkLocalIPv6Address": "",
"LinkLocalIPv6PrefixLen": 0,
"SecondaryIPAddresses": null,
"SecondaryIPv6Addresses": null,
"EndpointID": "17c85035c3bbdac8491c95ca356f68e1c3f183fe2c0a475a0131649a255f4b20",
"Gateway": "172.17.0.1",
"GlobalIPv6Address": "",
"GlobalIPv6PrefixLen": 0,
"IPAddress": "172.17.0.2",
"IPPrefixLen": 16,
"IPv6Gateway": "",
"MacAddress": "56:84:7a:41:d4:00"
```

```
b8b05d1340023fdde3b1f2ad5425d16",
lc3f183fe2c0a475a0131649a255f4b20",
```

```
    "GlobalIPv6Address": null,
    "GlobalIPv6PrefixLen": 0,
    "IPAddress": "172.17.0.2",
    "IPPrefixLen": 16,
    "IPv6Gateway": "",
    "MacAddress": "56:84:7a:41:d4:00"
}
```

# Docker networking

---

- Docker has its own NIC (docker0) the same subnet
- Used to connect to containers from the linux host

```
osc-guy-van-eec@lnx-guy-van-eec:~$ ip a s
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 00:0c:29:3a:c7:22 brd ff:ff:ff:ff:ff:ff
    altname enp2s1
    altname enx000c293ac722
    inet 192.168.204.128/24 brd 192.168.204.255 scope global dynamic noprefixroute ens33
        valid_lft 1150sec preferred_lft 1150sec
    inet6 fe80::20c:29ff:fe3a:c722/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
3: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether ca:d9:ff:de:72:bc brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::c8d9:ffff:fedc:72bc/64 scope link proto kernel ll
        valid_lft forever preferred_lft forever
```

# Docker networking

---

- You can check your access to your containers from your linux vm

```
osc-guy-van-eec@lnx-guy-van-eec:~$ ping 172.17.0.2
PING 172.17.0.2 (172.17.0.2) 56(84) bytes of data.
64 bytes from 172.17.0.2: icmp_seq=1 ttl=64 time=0.272 ms
64 bytes from 172.17.0.2: icmp_seq=2 ttl=64 time=0.107 ms
64 bytes from 172.17.0.2: icmp_seq=3 ttl=64 time=0.119 ms
^C
--- 172.17.0.2 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2051ms
rtt min/avg/max/mdev = 0.107/0.166/0.272/0.075 ms
osc-guy-van-eec@lnx-guy-van-eec:~$ |
```

- E.g.: webserver on container (port 80)

```
osc-guy-van-eec@lnx-guy-van-eec:~$ curl http://172.17.0.2/
<html><body><h1>It works!</h1></body></html>
osc-guy-van-eec@lnx-guy-van-eec:~$
```

# Docker networking

---

- Bridge network is only available from the host itself !
- Services should be made available outwards via port forwarding
  - **docker run -d -p 80:80 --name my\_web <container image name>**
  - E.g. Apache webserver: **docker run -d -p 8080:80 httpd**



It works!

# Docker networking

---

- BEWARE ! You can only allocate a port on your lnx host ONCE
- So if you have multiple http containers / daemons => Allocate to other ‘external’ port:

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker run -d -p 80:80 httpd
08086d54407a2b424723a1fd456a207a5fa5b1dbd9c4b7e284ee7a9cd08dbed3
docker: Error response from daemon: failed to set up container networking: driver failed programming external connectivity on endpoint distracted_wright (3a1fe1c6ed3ee1145a07f89034fa0118983da229
d479f0b97060947e404b69a2): failed to bind host port for 0.0.0.0:80:172.17.0.3:80/tcp: address already in use

Run 'docker run --help' for more information
osc-guy-van-eec@lnx-guy-van-eec:~$ docker run -d -p 8081:80 httpd
050b8e7b59d2a0bf2439c1a85369cdb7161dbb06b23021746353f48efab91e89
osc-guy-van-eec@lnx-guy-van-eec:~$ |
```

# Docker networking

---

- Standard install uses iptables to control access
  - **sudo iptables -L**

```
Chain DOCKER (1 references)
target      prot opt source          destination
ACCEPT      tcp   --  anywhere        172.17.0.4      tcp  dpt:http
ACCEPT      tcp   --  anywhere        172.17.0.3      tcp  dpt:http
ACCEPT      tcp   --  anywhere        172.17.0.2      tcp  dpt:http
DROP       all   --  anywhere        anywhere
```

# Docker logs

---

- Starting a container in foreground will put all logs on the terminal

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker run --name my_nginx nginx:1.27.0
Unable to find image 'nginx:1.27.0' locally
1.27.0: Pulling from library/nginx
e4fff0779e6d: Pull complete
9c4e4e507b08: Pull complete
85664e7b25d6: Pull complete
e9d00ba15cc0: Pull complete
ca3474dd2ca6: Pull complete
472851929ad7: Pull complete
ed60582f2661: Pull complete
Digest: sha256:98f8ec75657d21b924fe4f69b6b9bff2f6550ea48838af479d8894a852000e40
Status: Downloaded newer image for nginx:1.27.0
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/11/08 10:15:19 [notice] 1#1: using the "epoll" event method
2025/11/08 10:15:19 [notice] 1#1: nginx/1.27.0
2025/11/08 10:15:19 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/11/08 10:15:19 [notice] 1#1: OS: Linux 6.12.48+deb13-amd64
2025/11/08 10:15:19 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1073741816:1073741816
2025/11/08 10:15:19 [notice] 1#1: start worker processes
2025/11/08 10:15:19 [notice] 1#1: start worker process 30
2025/11/08 10:15:19 [notice] 1#1: start worker process 31
```

# Docker logs

- Starting a container in detached mode  
(-d) No longer shows logs !
- Get logs by using **docker logs <container-id or naam>**

```
osc-guy-van-eec@lnx-guy-van-eec:~$ docker logs my_nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-te
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-pr
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/11/08 10:15:19 [notice] 1#1: using the "epoll" event method
2025/11/08 10:15:19 [notice] 1#1: nginx/1.27.0
2025/11/08 10:15:19 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14
2025/11/08 10:15:19 [notice] 1#1: OS: Linux 6.12.48+deb13-amd64
2025/11/08 10:15:19 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1073741816:
2025/11/08 10:15:19 [notice] 1#1: start worker processes
2025/11/08 10:15:19 [notice] 1#1: start worker process 30
2025/11/08 10:15:19 [notice] 1#1: start worker process 31
2025/11/08 10:16:27 [notice] 1#1: signal 2 (SIGINT) received, exiting
2025/11/08 10:16:27 [notice] 30#30: exiting
2025/11/08 10:16:27 [notice] 30#30: exit
2025/11/08 10:16:27 [notice] 31#31: exiting
2025/11/08 10:16:27 [notice] 31#31: exit
2025/11/08 10:16:27 [notice] 1#1: signal 17 (SIGCHLD) received from 30
2025/11/08 10:16:27 [notice] 1#1: worker process 30 exited with code 0
2025/11/08 10:16:27 [notice] 1#1: signal 29 (SIGIO) received
2025/11/08 10:16:27 [notice] 1#1: signal 17 (SIGCHLD) received from 31
2025/11/08 10:16:27 [notice] 1#1: worker process 31 exited with code 0
2025/11/08 10:16:27 [notice] 1#1: exit
osc-guy-van-eec@lnx-guy-van-eec:~$
```

# Docker logs

---

- Stored in a separate file per container
- Location: /var/lib/docker/containers/<container-id>/<container-id>-json.log

```
osc-guy-van-eec@lnx-guy-van-eec:~$ sudo cat /var/lib/docker/containers/050b8e7b59d2a0bf2439c1a85369cdb7161dbb06b23021746353f48efab91e89/050b8e7b59d2a0bf2439c1a85369cdb7161dbb06b23021746353f48efab91e89-json.log
{"log": "AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using 172.17.0.3. Set the 'ServerName' directive globally to suppress this message\n", "stream": "stderr", "time": "2025-11-08T10:10:10.840198384Z"}
{"log": "AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using 172.17.0.3. Set the 'ServerName' directive globally to suppress this message\n", "stream": "stderr", "time": "2025-11-08T10:10:10.846912295Z"}
{"log": "[Sat Nov 08 10:10:10.852052 2025] [mpm_event:notice] [pid 1:tid 1] AH00489: Apache/2.4.65 (Unix) configured -- resuming normal operations\n", "stream": "stderr", "time": "2025-11-08T10:10:10.852514562Z"}
{"log": "[Sat Nov 08 10:10:10.852161 2025] [core:notice] [pid 1:tid 1] AH00094: Command line: 'httpd -D FOREGROUND'\n", "stream": "stderr", "time": "2025-11-08T10:10:10.852531132Z"}
osc-guy-van-eec@lnx-guy-van-eec:~$
```

# Today's LAB

---

- 1) Install docker on your Inx- VM THE RIGHT WAY
- 2) Docker hello-world (yay ! Hello world)
- 3) Observe docker container and container image behavior
- 4) Dabble with tags
- 5) Check out the docker logs
- 6) Explore Docker Networking